

# Linux File I/O : Reading and Writing

**Advanced Embedded Linux  
Development  
with Dan Walkes**



University of Colorado **Boulder**

**Learning objectives:**

**Reading and Writing Files**

**How umask influences file  
permissions**

**Introduction to strace**

# Reading file content

- `read()`
  - Returns:
    - Number of bytes read, or -1 on error (and sets `errno`)
    - Return values < `count` mean no more bytes are available now (end of file, interrupted by signal)
- READ(2)** Linux Programmer's Manual

**NAME** read - read from a file descriptor

**SYNOPSIS**

```
#include <unistd.h>

ssize_t read(int fd, void *buf, size_t count);
```

**DESCRIPTION**

read() attempts to read up to `count` bytes from file descriptor `fd` into the buffer starting at `buf`.

# Read File Simple Example

```
int main(int argc, char **argv)
{
    int rc=1;
    if( argc != 2 ) {
        printf("Usage: %s filename\n",argv[0]);
        printf("    The file specified as \"filename\" will be created with zero size if it does not exist.\n");
    } else {
        const char *filename = argv[1];
        int fd=open(filename,
                    O_RDWR|O_CREAT,
                    S_IRWXU|S_IRWXG|S_IRWXO);
        char buff[1];
        if( fd == -1 ) {
            printf("Error opening %s\n",filename);
        } else {
            int readlen = read(fd,buff,1);
            rc = 0;
            if( readlen == -1 ) {
                printf("Error with read(), errno is %d (%s)\n",errno,strerror(errno));
                rc = -1;
            }
            if( readlen != 0 ) {
                printf("read %d bytes, first byte was %c\n",readlen,buff[0]);
            }
            else {
                printf("Read 0 bytes (empty file)\n");
            }
            close(fd);
        }
    }
    return rc;
}
```

- Read the first byte of a file
- Create the file with zero size if it doesn't exist.

# Create File and Umask

```
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ ./run-empty-file-read.sh  
Read 0 bytes (empty file)  
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ ls -l  
total 24  
-wxrwxr-x 1 aesd aesd 0 Jan 26 18:06 emptyfile  
-rwxrwxr-x 1 aesd aesd 8624 Jan 26 18:17 readfile  
-rw-rw-r-- 1 aesd aesd 1139 Jan 26 18:14 readfile.c  
-rwxrwxr-x 1 aesd aesd 59 Jan 26 18:14 run-empty-file-read.sh  
-rwxrwxr-x 1 aesd aesd 181 Jan 26 18:14 run-pipe-read.sh
```

- RWX for others?

```
int fd=open(filename,  
            O_RDWR|O_CREAT,  
            S_IRWXU|S_IRWXG|S_IROWXO);
```

- What happened to the write permission?
  - Cleared by umask
- What is umask?
  - Per-process way to restrict permissions set on created files. (& ~umask)

```
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ umask  
0002
```

# Read File Simple Example

- What happens if I change my umask to 0003?

```
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ umask 0003  
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ umask  
0003  
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ rm emptyfile  
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ ./run-empty-file-read.sh  
Read 0 bytes (empty file)
```

```
int fd=open(filename,  
            O_RDWR|O_CREAT,  
            S_IRWXU|S_IRWXG|S_IRWXO);
```

```
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ ls -l  
total 24  
-rwxrwxr-- 1 aesd aesd 0 Jan 26 18:22 emptyfile  
-rwxrwxr-- 1 aesd aesd 8624 Jan 26 18:22 readfile  
-rw-rw-r-- 1 aesd aesd 1139 Jan 26 18:14 readfile.c  
-rwxrwxr-x 1 aesd aesd 59 Jan 26 18:14 run-empty-file-read.sh  
-rwxrwxr-x 1 aesd aesd 181 Jan 26 18:14 run-pipe-read.sh
```

- Now both w and x are cleared.

# System Call Logging with strace

```
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ strace ./readfile emptyfile
execve("./readfile", ["./readfile", "emptyfile"], 0x7ffe16ee9088 /* 22 vars */) = 0
brk(NULL)                                     = 0x563b0f645000

write(1, "Opening emptyfile in blocking mo"..., 35Opening emptyfile in blocking mode
) = 35
openat(AT_FDCWD, "emptyfile", O_RDWR|O_CREAT, 0777) = 3
write(1, "Reading 1 byte from emptyfile\n", 30Reading 1 byte from emptyfile
) = 30
read(3, "", 1)                                = 0
write(1, "Read 0 bytes (empty file)\n", 26Read 0 bytes (empty file)
) = 26
close(3)                                       = 0
exit_group(0)
```

- strace is a handy tool for debugging application interaction with the kernel through syscall
- No source code required!



# Non-Blocking Reads

```
int fd=open(filename,
            O_RDONLY|
            /**
             * When opening in blocking mode, use O_CREAT|O_RDWR to support creating the file if it doesn't exist
             * and to allow opening pipes on Linux when no readers are available (see http://man7.org/linux/man-pages/man7/fifo.7.html)
             * blocking mode.
            */
            block ? O_CREAT|O_RDWR : O_NONBLOCK,
            S_IRWXU|S_IRWXG|S_IRWXO);
```

Use O\_NONBLOCK on open to open in non blocking mode

```
int readlen = read(fd,buf,1);
rc = 0;
if( readlen == -1 ) {
    if( !block && errno == EAGAIN ) {
        printf("EAGAIN returned from read(), no data available to read\n");
    } else {
        printf("Error returned from read(), errno is %d (%s)\n",errno,strerror(errno));
    }
}
```

check for EAGAIN to see if data is available

- **O\_NONBLOCK** allows your thread needs to do work while waiting for fifo/pipe data.

# Non-Blocking Reads Example

```
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ ./run-pipe-read.sh
make: Nothing to be done for 'all'.
Attempting read on a pipe, will block
Run "echo 'hi' > pipe" from another window to unblock
Opening pipe in blocking mode
Reading 1 byte from pipe
echo 'hi' > pipe
Read 1 bytes, first byte was h
```

```
aesd@aesd-VirtualBox:~/aesd-lectures/lecture5$ ./run-pipe-read-nonblocking.sh
make: Nothing to be done for 'all'.
Attempting read on a pipe in non-blocking mode

Using nonblocking mode as specified by argument
Opening pipe in nonblocking mode
Reading 1 byte from pipe
EAGAIN returned from read(), no data available to read
```

```
#!/bin/sh
make
rm pipe
mkfifo pipe
echo "Attempting read on a pipe, will block"
echo "Run \"echo 'hi' > pipe\" from another window to unblock"
./readfile pipe
```

```
#!/bin/sh
make
rm -f pipe
mkfifo pipe
# Open a command for reading the pipe and put in the background
# This prevents the next command from blocking
# See https://unix.stackexchange.com/questions/366219/pr
echo < pipe &
# Open a command to write to the pipe, but don't write a
# EOF
exec 3>pipe
echo "Attempting read on a pipe in non-blocking mode"
./readfile pipe nonblock
```

# Writing Files

- `write()`

```
WRITE(2)                               Linux Programmer's Manual

NAME
    write - write to a file descriptor

SYNOPSIS
    #include <unistd.h>

    ssize_t write(int fd, const void *buf, size_t count);

DESCRIPTION
    write() writes up to count bytes from the buffer starting at buf to the file referred to by the file descriptor fd.
```



# Synchronized I/O

- Is data written to disk when write() completes successfully?
  - No - the write is into a kernel buffer.
- Why?
  - Disks are slow!



# Synchronized I/O

- When is this a problem?
  - Unclean shutdown
- How can we manage the risk?
  - `fsync()`, `fdatasync()`, `O_SYNC`
  - Configure disk write caching.
- Avoid `sync()` for performance reasons